

TLO: Topology-Lattice Obfuscation for Smart Contracts

Anonymous Submission

Abstract

We present TLO (Topology-Lattice Obfuscation), a practical on-chain compiler that produces publicly evaluable locked circuits whose internal gate-control representation is LWE-hidden and whose topology is engineered to defeat a defined suite of structural reverse-engineering attacks. Security derives from a two-layer defense: a topology layer using structural mixing that empirically defeats structural and statistical attacks, and an LWE layer that hides the gate-control representation via on-chain inner products.

The LWE layer provides representation hiding, not semantic security; unlocking security reduces to secret search in the input space (and/or breaking LWE). Attackers *can* simulate evaluation offline for any input x by computing $s(x) = H(x)$. Our security claims are heuristic and relative to our defined attack evaluation matrix; we achieve 6/6 resistance within that framework at $\sim 2.58M$ gas (8.6% of block limit). We give uniform-secret LWE security estimates (~ 49 -bit classical with $n=64$; see §5.3), providing post-quantum representation hiding (assuming LWE quantum-resistance). Target applications include predicates with eventually-expiring secrets (honeypots, sealed-bid auctions, lotteries).

1 Introduction

Smart contracts are fully transparent. Anyone can read bytecode, analyze logic, and exploit vulnerabilities. This conflicts with applications requiring hidden logic: cryptographic honeypots, MEV-resistant execution, sealed-bid auctions, and private liquidation thresholds.

Indistinguishability obfuscation (iO) provides strong security but requires impractical overhead ($10^6 \times$). We take a different approach: *heuristic obfuscation* in the tradition of local mixing [1], providing representation hiding and attack-surface hardening through two complementary layers:

Two-Layer Security Model:

1. **Layer 1 (Topology):** Structural mixing empirically defeats structural/statistical attacks. *Security: heuristic.*
2. **Layer 2 (LWE):** On-chain inner products hide gate-control representation. *Security: computational (representation hiding under ~ 49 -bit uniform-secret LWE).*

Key property: Unlocking reduces to secret search, not to preventing offline evaluation.

Terminology note: We use “obfuscation” informally, in the tradition of code obfuscation and heuristic obfuscation (e.g., Canetti et al. [1]), to denote representation hiding and attack-surface hardening. We do *not* claim indistinguishability obfuscation, virtual-black-box security, or any standard strong obfuscation definition.

1.1 Contributions

1. **TLO Framework:** Two-layer heuristic obfuscation combining topology mixing with on-chain LWE inner products ($\sim 2.58M$ gas for $n=64$).
2. **Structural Mixing:** Wire selection defeating structural/statistical attacks (heuristic, empirically validated).
3. **Wrong-Key-Garbage:** Attackers *can* simulate with arbitrary keys, but incorrect keys yield garbage. This does not improve brute-force search over the secret input space; it hides direct access to the gate-control representation.
4. **On-Chain LWE Representation Hiding:** Control functions hidden via LWE ciphertexts with full inner product computation.
5. **Post-Quantum Representation Hiding:** ~ 49 -bit classical / ~ 45 -bit quantum estimate via uniform-secret LWE ($n=64$); ~ 81 -bit for $n=128$; see §5.3.

1.2 Scope

We claim: Representation hiding based on LWE hardness; topology heuristics empirically validated;

6/6 attack resistance *within our defined matrix*; post-quantum representation hiding.

We do NOT claim: iO security; VBB security; semantic hiding beyond black-box behavior; prevention of offline evaluation.

2 Preliminaries

2.1 Learning With Errors

Definition 2.1 (LWE [6]). For dimension n , modulus q , and error distribution χ , the LWE problem is: given $(A, As + e \bmod q)$ where $A \in \mathbb{Z}_q^{m \times n}$, $s \in \mathbb{Z}_q^n$, $e \leftarrow \chi^m$, distinguish from uniform (A, u) .

LWE is believed quantum-resistant and forms the basis for post-quantum cryptography standards (ML-KEM) [2]. Our parameters ($n=64$, $q=65521$) are smaller than NIST profiles; see §5.3 for security estimates.

2.2 LWE Control Function Hiding

We hide each gate’s control function via LWE ciphertexts. Each CF bit is encoded as (a, b) where $b = \langle a, s_{\text{enc}} \rangle + e + \text{bit} \cdot q/2$. At encryption time, $s_{\text{enc}} = H(\text{secret})$. At evaluation time, the contract derives $s(x) = H(x)$ from the candidate input—decryption succeeds iff $x = \text{secret}$.

Important: Since $s(x) = H(x)$ is publicly computable, anyone can evaluate the locked predicate offline for arbitrary inputs. The LWE layer hides the *representation* (CF bits), not the *functionality*.

2.3 Reversible Circuits

Definition 2.2 (Reversible Gate). A gate $g = (a, c_1, c_2, c_f)$ operates on n wires: active wire a is XORed with $c_f(c_1, c_2)$ where $c_f : \{0, 1\}^2 \rightarrow \{0, 1\}$ is one of 16 control functions.

Gates are self-inverse: $g(g(s)) = s$. This enables commit-reveal protocols where the solver demonstrates knowledge without revealing the secret.

3 The Topology Layer

The topology layer is a reversible circuit mixing design that empirically defeats structural and statistical attacks through wire selection, without cryptographic primitives.

3.1 Wire Selection Algorithm

Structural mixing selects wires to defeat pattern detection:

1. **Non-pow2 distances:** Control wires at distances $d \notin \{1, 2, 4, 8, \dots\}$ from active wire defeats butterfly/FFT pattern detection

2. **Uniform wire usage:** Prefer underused wires; defeats chi-squared statistical attacks
3. **Irregular layers:** Varying gates per layer (e.g., 30–70 for 256 wires) defeats regularity detection
4. **64+ wires:** Sufficient width defeats diagonal correlation (Pearson $r < 0.10$)

Listing 1: Wire Selection (pseudocode)

```
def select_control_wire(active, usage, target):
    # Choose non-pow2 distance
    d = random_choice(non_pow2_distances)
    candidate = (active + d) % num_wires

    # Prefer underused wires (70% prob)
    if random() < 0.7 and usage[candidate] < target:
        return candidate

    # Otherwise, find underused alternative
    return find_underused_wire(usage, target)
```

3.2 Topology Attack Resistance

Attack	Type	Defense	Mechanism
Compression	Structural	Topology	No duplicate gates
PatternMatch	Structural	Topology	Random CF cycling
Structural	Structural	Topology	Non-pow2 distances
Statistical	Statistical	Topology	Uniform wire usage
DiagCorrelation	Statistical	Topology	64+ wires

Table 1: Topology empirically defeats structural/statistical attacks in our evaluation.

Key insight: Unlike butterfly or derangement topologies that only rearrange gates, structural mixing has anti-attack properties *built into wire selection*.

4 LWE for Representation Hiding

4.1 The RainbowTable Problem

RainbowTable is a *semantic* attack—it matches truth-table behavior, not structure:

1. Extract subcircuit from obfuscated circuit
2. Evaluate subcircuit on sample inputs
3. Match behavior against pre-computed lookup table

Topology cannot defeat this. Any structural transformation preserves semantic behavior of reversible circuits.

4.2 What LWE Does and Does Not Provide

Critical clarification: Attackers can compute $s(x) = H(x)$ for any input x and thus can freely simulate the contract’s 1-bit output offline. Therefore, our

LWE layer does *not* prevent offline evaluation of the locked predicate.

Instead, it hides the internal gate-control representation and aims to frustrate semantic *reverse-engineering* attacks such as RainbowTable-style structural matching.

On-Chain Inner Product:	Control functions encoded as LWE ciphertexts (a, b) . At encryption: $s_{\text{enc}} = H(\text{secret})$. At evaluation: $s(x) = H(x)$.
What this provides:	Representation hiding—CF bits are not directly readable from the static artifact without choosing an input x and computing $H(x)$.
What this does NOT provide:	Prevention of offline evaluation. Unlocking difficulty is governed by secret entropy.

Proposition 4.1 (LWE-Based Representation Hiding). *Assuming the hardness of LWE, given only the obfuscated circuit, no PPT adversary can recover a non-negligible fraction of the gate control-function bits c_f with non-negligible advantage.*

Proof sketch: Each control-function bit is encoded as an LWE ciphertext (a, b) under a secret $s_{\text{enc}} = H(\text{secret})$. Recovering c_f from these ciphertexts without knowledge of s_{enc} reduces to distinguishing LWE samples from uniform. \square

5 Security Analysis

5.1 Two-Layer Security Model

TLO provides security through complementary layers with different bases:

1. **Topology layer (heuristic):** Empirically defeats structural/statistical attacks through wire selection. *Empirically validated, not proven.*
2. **LWE layer (computational):** Provides representation hiding via on-chain inner products. *Based on LWE hardness (~ 49 -bit with $n=64$; see §5.3).*
3. **Unlocking security:** Reduces to secret search in the input space (and/or breaking LWE). Attackers can evaluate the locked predicate offline for arbitrary inputs.

Definition 5.1 (Attack-Matrix Resistance). Given a fixed finite set of attack classes \mathcal{A} and an evaluation metric, we say that an obfuscation scheme achieves \mathcal{A} -resistance if no PPT adversary implementing any attack in \mathcal{A} succeeds with probability exceeding a specified threshold (here, empirically negligible on our benchmarks).

Theorem 5.2 (TLO Attack-Matrix Resistance). *Under our LWE-based representation-hiding assumption, our empirical topology heuristics, and the wrong-key-gives-garbage property, TLO achieves \mathcal{A} -resistance for our 6-class attack matrix on the evaluated parameter sets.*

Proof (informal): Structural/statistical attacks in \mathcal{A} are thwarted by the topology layer in our experiments. RainbowTable-style attacks in \mathcal{A} that rely on explicit control-function recovery are blocked by LWE-based representation hiding. We do not claim resistance to all possible attacks outside \mathcal{A} . \square

5.2 Attack Evaluation Matrix

Attack	Defense	Basis	Status
Compression	Topology	Structural	BLOCKED
PatternMatch	Topology	Structural	BLOCKED
Structural	Topology	Structural	BLOCKED
Statistical	Topology	Statistical	BLOCKED
DiagCorrelation	Topology	Statistical	BLOCKED
RainbowTable	LWE	Semantic	BLOCKED*

Table 2: TLO attack resistance within our defined matrix. *Blocks structural rainbow-table matching; does not prevent black-box evaluation.

5.3 Security Estimates

Uniform-Secret LWE. TLO derives the LWE secret as $s_{\text{enc}} = H(\text{secret})$, producing a *uniform* secret over \mathbb{Z}_q^n rather than a small-coefficient secret. This variant is *harder* to attack: primal (uSVP) attacks fail when $\|s\| \approx \sqrt{n} \cdot q/2$. We validated this via BKZ attacks using fpyll—BKZ-50 on $n=16$ failed after 200+ iterations.

Using dual-attack analysis (which applies regardless of secret distribution), our parameters ($n=64$, $q=65521$, $\sigma=\sqrt{q}/4$, $m=2560$ samples) yield ~ 49 -bit security. This is suitable for *eventually-expiring secrets* with hour-to-day lifetimes. For $n=128$: ~ 81 -bit; for $n=256$: ~ 132 -bit (NIST-level).

Hash-Compare Baseline: A simple $H(\text{secret}) == H(\text{input})$ check costs $\sim 45K$ gas but provides *no* representation hiding—the predicate structure is visible on-chain.

Multi-Bit Output: The Key Distinction. Hash-compare returns a 1-bit output (true/false). TLO circuits compute an N -bit output that can encode hidden parameters, computed results, or payloads revealed only on correct input. Both implement *point functions*—predicates meaningful only at $x = \text{secret}$ —but TLO provides a hidden payload, not just confirmation.

Approach	Output	What's Hidden
Hash-compare	1 bit	Secret value only
TLO	N bits	Secret + hidden computation

Comparison: In both designs, an attacker can evaluate the predicate on arbitrary inputs (on-chain or off-chain), so the difficulty of finding a satisfying input before expiry primarily depends on the secret’s entropy and application-level constraints, not on LWE. TLO’s advantage over hash-compare is: (1) multi-bit hidden output, and (2) hiding the internal predicate representation (topology + control functions). The $57\times$ gas premium buys multi-bit hidden computation, not stronger unlocking security.

5.4 Post-Quantum Security

TLO’s LWE layer provides post-quantum *representation hiding* (assuming LWE quantum-resistance):

- **Topology layer:** No cryptographic assumptions
- **Lattice layer:** LWE is believed quantum-resistant

5.5 Assumptions

1. **LWE hardness:** Learning With Errors is computationally hard.
2. **Topology empirical security:** Wire selection defeats structural attacks in our evaluation (heuristic, not proven).
3. **Contract correctness:** Expiry logic is correctly implemented.

6 Implementation

6.1 Contract Architecture

TLO requires no external infrastructure—just a standard smart contract with timestamp-based expiry:

Listing 2: TLOHoneypot Contract

```
contract TLOHoneypot {
    uint256 public secretExpiry;
    bytes32 public commitHash;

    function check(bytes32 s) external view
        returns (bool) {
        require(block.timestamp < secretExpiry);
        return evaluate(s);
    }

    function commit(bytes32 h) external {
        commitHash = h;
    }

    function reveal(bytes32 s) external {
        require(keccak256(abi.encode(s,
            msg.sender)) == commitHash);
        require(evaluate(s));
        // Transfer reward
    }
}
```

Deployment: Set `secretExpiry` at deployment.

6.2 Gas Costs

Measured on 64-wire/640-gate circuits (Tenderly-confirmed):

LWE n	Security	Gas	Block %
16	~22-bit	744K	2.5%
32	~22-bit	1.27M	4.2%
64	~49-bit	2.58M	8.6%
128	~81-bit	4.9M	16.3%

Table 3: Gas costs by LWE dimension (uniform-secret security estimates for representation hiding).

7 Evaluation

7.1 Attack Resistance

We evaluated TLO against 14 attack implementations across 6 attack classes. All configurations achieve 6/6 resistance within our defined attack matrix:

LWE n	Score	Gas	Security
16	6/6	744K	~22-bit
32	6/6	1.27M	~22-bit
64	6/6	2.58M	~49-bit
128	6/6	4.9M	~81-bit

7.2 Comparison with Alternatives

Property	TLO	Hash	iO	TEE
Attack matrix	6/6	0/6	6/6	6/6
Repr. hiding	Yes	No	Yes	Yes
Unlocking	Search	Search	Search	Key
Gas (check)	2.58M	45K	$10^6 \times$	$1 \times$
Infrastructure	None	None	None	HW
Post-quantum	Yes	Yes	Depends	No
Practical	Yes	Yes	No	Yes

8 Applications

8.1 Valid Applications

TLO is designed for predicates with *eventually-expiring* secrets where representation hiding provides value:

- **Cryptographic honeypots:** Reward condition is burned once triggered
- **Sealed-bid auctions:** Bids revealed at settlement
- **Lotteries/prediction markets:** Outcomes revealed after close
- **MEV protection:** Order flow is short-lived
- **Dark pools:** Trade conditions expire quickly

8.2 Invalid Applications

TLO is *not* intended for long-lived static secrets or applications requiring semantic security:

- Long-term decryption keys
- Permanent signing keys
- Static liquidation thresholds (see §9)

9 Limitations

Theoretical: Topology security is empirical (heuristic, not proven). We do not claim iO-level indistinguishability or semantic security.

Practical: TLO with $n=64$ LWE requires $\sim 2.58M$ gas (8.6% of block limit). Lower security configurations available for cost-sensitive applications.

What TLO does NOT provide:

- **Prevention of offline evaluation:** Attackers can evaluate the locked predicate for arbitrary inputs
- **Indistinguishability:** Two circuits have distinguishable obfuscations
- **Universal security:** Only resists our 6 attack classes
- **Forward secrecy:** Expired secrets may be analyzed retroactively
- **Security beyond hash-compare for unlocking:** Unlocking difficulty is governed by secret entropy, not LWE

10 Related Work

Indistinguishability Obfuscation: Theoretical iO [3, 5] provides strong security but requires impractical overhead.

Local Mixing: Canetti et al. [1] explore heuristic obfuscation via local perturbations of reversible circuits. Like their work, we provide candidate constructions with empirical (not proven) security.

Compute-and-Compare: Goyal-Koppula-Waters [4] and Wichs-Zirdelis [7] introduced C&C for evasive functions. We apply similar ideas to control function hiding.

Smart Contract Privacy: Previous work uses ZK-SNARKs (Tornado Cash) or TEEs (Secret Network). TLO provides a new point in the design space: on-chain representation hiding without trusted hardware.

11 Conclusion

TLO provides practical heuristic obfuscation for smart contracts through two-layer defense: a topology layer (heuristic) empirically defeats structural/statistical attacks, while on-chain LWE inner products provide representation hiding for control functions.

TLO achieves 6/6 resistance within our defined attack matrix at $\sim 2.58M$ gas ($n=64$, ~ 49 -bit uniform-secret LWE security for representation hiding). Post-quantum representation hiding (assuming LWE quantum-resistance). Deployment requires only a standard smart contract with timestamp expiry.

Key contributions: On-chain LWE inner products for representation hiding; honest characterization of security properties; discovery that uniform-secret LWE provides stronger security than small-secret variants.

Honest scope: Unlocking security reduces to secret search, not to preventing offline evaluation. TLO hides *how* the predicate is implemented, not *whether* it can be evaluated.

Code and attack suite: <https://github.com/igor53627/tlo>

References

- [1] R. Canetti, C. Chamon, E. Mucciolo, and A. Rückenstein. Towards general-purpose program obfuscation via local mixing. Cryptology ePrint Archive, Report 2024/006, 2024.
- [2] National Institute of Standards and Technology. Module-Lattice-Based Key-Encapsulation Mechanism Standard (ML-KEM). NIST FIPS 203, 2024.
- [3] S. Garg, C. Gentry, S. Halevi, M. Raykova, A. Sahai, and B. Waters. Candidate indistinguishability obfuscation and functional encryption for all circuits. In *FOCS*, 2013.
- [4] R. Goyal, V. Koppula, and B. Waters. Lockable obfuscation. In *FOCS*, 2017.
- [5] A. Jain, H. Lin, and A. Sahai. Indistinguishability obfuscation from well-founded assumptions. In *STOC*, 2021.
- [6] O. Regev. On lattices, learning with errors, random linear codes, and cryptography. In *STOC*, 2005.
- [7] D. Wichs and G. Zeldelis. Obfuscating compute-and-compare programs under LWE. In *FOCS*, 2017.
- [8] M. R. Albrecht, R. Player, and S. Scott. On the concrete hardness of Learning with Errors. *Journal of Mathematical Cryptology*, 9(3):169–203, 2015.